

1 Theoretical Questions

1. AI is a broad field with many interpretations, and there are thus numerous ways to define this field. One broad definition is that AI is the study and construction of agents that do the right thing. Though what is considered the right thing is also up for debate. An AI should be able to make decisions or act in a way that could mimic human intelligence, thus it should emit intelligent behavior. Moreover, what we may consider intelligent behavior is possible to define in many ways. One way to define it is that an AI should try to maximize expected utility. This means that the AI should try to achieve some goal within a given set of constraints in a way that maximizes the expected outcome (Russell and Norvig 2021, pp. 7–22).
2. The Turing test is an experiment created by Alan Turing to see if a machine could exhibit intelligent enough behaviour to such an extent that a human interacting with it, for instance posing and answering questions, could not distinguish it from another human. A new proposal of the Turing test includes interaction not only through answering and posing questions, but actual physical interaction with humans and objects in the real world, which would require for instance computer vision and robotics in addition to natural language processing (Russell and Norvig 2021, p. 23).
3. Rationality is defined as doing the right thing. Thinking rationally can be described as the ability to make predictions about the future based on a sound understanding of how the world works. Acting rationally, however, is the ability to make decisions that maximize expected utility. An AI doesn't need to be able to perform rational thinking the way we humans do to act rationally. It should make decisions that are optimal given a set of goals or constraints, and may act rationally without rational thought if it for example has learned from past runs or follows a set of rules (Russell and Norvig 2021, pp. 19–22).
4. According to Aristotle, there is a connection between knowledge and action due to the fact that to reach a goal, one has to have knowledge about the outcomes of an action, and thus what actions are required to reach a specific goal. Aristotle himself proposed an algorithm which explains how his argument can be used to implement the idea in AI. One would assume the goal state and then work backwards to find a sequence of actions that would lead to that goal state, i.e. regression.
 - (a) Newell and Simon were the first AI researchers to implement Aristotle's ideas.
 - (b) The name of the program they developed was General Problem Solver, which is a planning system that uses a greedy algorithm together with regression to find a solution to a problem.(Russell and Norvig 2021, p. 25)
5.
 - The robot has a set of restricted actions, which include looking and going in four different directions. We can assume that the robot would look in a specific direction, e.g. forward, to see if there were any obstacles blocking it from reaching its goal, i.e. crossing the road. If there were, it would take a step left/right, and if there were none, simply go forward. It is after this action is taken that an elk suddenly crashes into the robot and smashes it. Based on this, the robot made decisions based on its available actions and the information it had at the time, and thus acted rationally. The robot could not have known that an elk would crash into it, and thus could not have avoided it.
 - The robot chose to cross the road on a green light, which from using its lookForward action, would mean that it was safe to cross the road. Again, the robot could not have known that a car would drive a red light and subsequently crash into it, and thus acted rationally. Even though it didn't reach its goal in any cases, it still acted within the actions and knowledge it had at its disposal to best reach its goal.
6. 6
7. 7
8. 8

Bibliography

Russell, Stuart and Peter Norvig (2021). *Artificial Intelligence: A Modern Approach, 4th Edition*. Pearson Education.